

# Libraries

```
In [9]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [11]: import pandas as pd
df = pd.read_csv('Diabetes.csv')

print(df)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..	...	...	...	...	...	...
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
In [12]: df.head()
```

Out[12]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeF
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

## DATA OVERVIEW

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

## DATA CLEANING

In [4]: df.isnull().sum()

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

In [5]: Duplicate Value

Duplicate values in df are: 0

In [5]: print ("Duplicate values in df are:" , df.duplicated().sum())

Duplicate values in df are: 0

```
In [ ]: Unique categories of Categorical Variables

In [6]: print(df['Outcome'].unique())

[1 0]
```

# Correlation


```
In [7]: df.corr()
```

Out[7]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000
BMI	0.017683	0.221071	0.281805	0.392573	0.197851
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185070
Age	0.544341	0.263514	0.239528	-0.113970	-0.042166
Outcome	0.221898	0.466581	0.065068	0.074752	0.130540

```
In [13]: plt.figure (figsize = [8,6], dpi = 130 )
plt.title ("Correlation Graph" , fontsize = 11 )
sns.heatmap (df.corr(), annot = True , cmap="coolwarm" )

Out[13]: <Axes: title={'center': 'Correlation Graph'}>
```



## Split X and y into training and testing sets

```
In [21]: X = pd.DataFrame (df , columns = ["Pregnancies" , "Glucose" , "BloodPressure" ,
                                           , "DiabetesPedigreeFunction" , "Age"]) # Features
y = df.Outcome # Target

X_train , X_test , y_train , y_test = train_test_split (X , y , test_size = 0.25)
```

```
In [22]: logreg = LogisticRegression (solver = "liblinear")

logreg.fit (X_train , y_train)
y_pred = logreg.predict(X_test)
y_predicted_proba = logreg.predict_proba(X_test)
```

```
In [23]: print ("Accuracy: " , metrics.accuracy_score (y_test , y_pred))
```

Accuracy: 0.8072916666666666

```
In [25]: confusion_matrix (y , logreg.predict (X))
```

```
Out[25]: array([[443, 57],
               [120, 148]], dtype=int64)
```

```
In [26]: cm = confusion_matrix (y , logreg.predict(X))

fig , ax = plt.subplots (figsize = (8,8))
ax.imshow (cm)
```

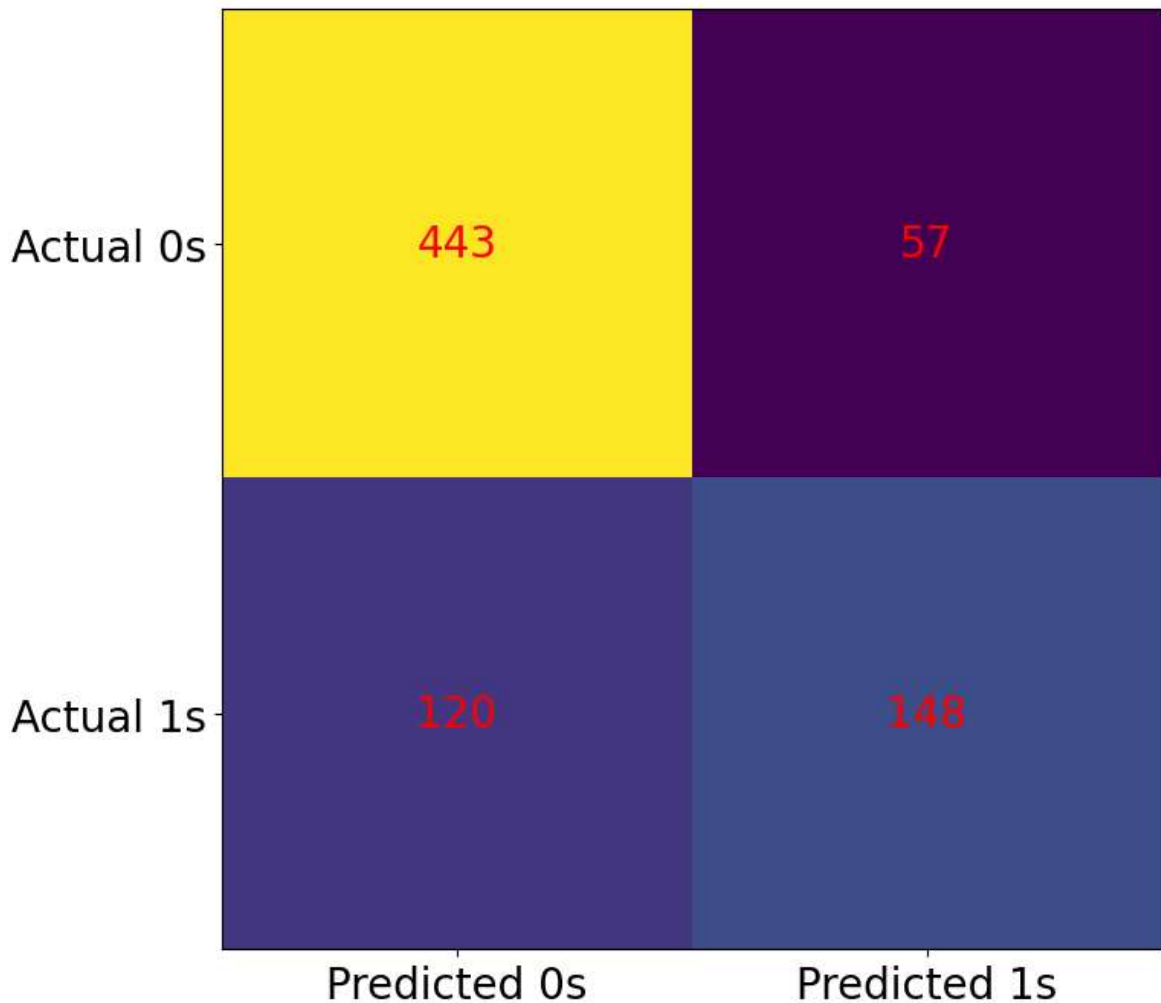
```

ax.grid (False)
ax.xaxis.set (ticks = (0 , 1) , ticklabels = ("Predicted 0s" , "Predicted 1s"))
ax.yaxis.set (ticks = (0 , 1) , ticklabels = ("Actual 0s" , "Actual 1s"))

ax.tick_params(axis='both', which='major', labels=20)
#ax.tick_params(axis='both', which='minor', labels=20)
#plt.xticks(fontsize=14, rotation=90)

ax.set_ylim (1.5 , -0.5)
for i in range (2):
    for j in range (2):
        ax.text (j , i , cm[i,j] , ha = "center" , va = "center" , color = "red",
plt.show()

```



```
In [27]: print (classification_report (y , logreg.predict (X)))
```

	precision	recall	f1-score	support
0	0.79	0.89	0.83	500
1	0.72	0.55	0.63	268
accuracy			0.77	768
macro avg	0.75	0.72	0.73	768
weighted avg	0.76	0.77	0.76	768

```
In [ ]:
```